

# Machine Learning A-Z

Obrusník Vít

CTU - Faculty of electrical engineering

**Abstract.** This document summarizes Machine Learning A-Z course available at Udemy. <https://www.udemy.com/machinelearning/learn/v4/overview>. Authors are Kirill Eremanko and Hadelin de Ponteves. All concepts are shown in Python and R. Course provides the students with code templates which are included in this document. Every section has a dataset to train and test the model.

## Table of Contents

Machine Learning A-Z .....	1
Introduction .....	2
1 Preprocessing .....	2
1.1 Problems with datasets .....	3
1.2 Code template in Python .....	3
1.3 Code template in R .....	3
2 Regression .....	5
2.1 Regression code template in Python .....	5
2.2 Regression code template in R .....	5
2.3 Simple Linear Regression .....	6
2.4 Multiple Linear Regression .....	8
Assumptions of Linear Regression .....	8
5 methods of building models .....	9
Backward elimination .....	9
Forward Selection .....	9
Bidirectional elimination .....	10
Code .....	10
2.5 Polynomial Regression .....	12
2.6 Support Vector Regression (SVR) .....	13
2.7 Decision Tree Regression .....	14
2.8 Random Forest Regression .....	15
2.9 Evaluating model performance .....	17
R-squared .....	17
Adjusted R-squared .....	17

3	Classification .....	18
3.1	Classification code template in Python .....	19
3.2	Classification code template in R .....	20
3.3	Logistic regression .....	21
3.4	K-Nearest neighbours .....	24
3.5	SVM (Support Vector Machine) .....	25
3.6	Kernel SVM .....	28
3.7	Naive Bayes .....	28
3.8	Decision Tree Classification .....	30
3.9	Random Forest Classification .....	32
3.10	Evaluating Classification model performance .....	33
	Confusion matrix .....	33
	Accuracy paradox .....	34
	CAP curve .....	35
4	Clustering .....	36
4.1	K-means clustering .....	37
	Elbow method to find optimal K .....	37
	Implementation .....	38
4.2	Hierarchical clustering .....	40
	Dendrograms .....	41
5	Association rule learning .....	43
5.1	Apriori .....	43
5.2	Eclat .....	46
6	Reinforcement learning .....	47
6.1	Upper Confidence Bound (UCB) .....	47
6.2	Thompson sampling .....	47
7	Natural language processing .....	48
8	Deep learning .....	49
8.1	Artificial Neural Network .....	49
8.2	Convolutional Neural Network .....	49
9	Dimensionality reduction .....	50
9.1	Principal Component Analysis (PCA) .....	50
9.2	Linear Discriminant Analysis (LDA) .....	50
9.3	Kernel PCA .....	50
10	Model selection and boosting .....	51
10.1	Model selection .....	51
10.2	XGBoost .....	51

## 1 Preprocessing

First of all we need to prepare our dataset. We can encounter several problems. Then we want to split the dataset.

## 1.1 Problems with datasets

### 1. Missing data

There might be some columns missing in the dataset. In that case our models won't work. We can either delete the incomplete rows or we can replace the missing data with average values.

### 2. Categorical data

Some parameters (e.g. City['Madrid', 'Barcelano', 'Sevilla']) need to be replaced with integer values (e.g. City[1, 2, 3]). We need to focus on not to fall into **Dummy Variable Trap**. More information: <http://www.algosome.com/articles/dummy-variable-trap-regression.html>

### 3. Splitting the dataset into training set and test set

We usually train the model on 80% of the dataset and test it on 20%.

### 4. Feature scaling

We might want to scale the data to have normalized magnitudes. This is not always necessary and depends on the method/library. All the methods that use Euclidian distance are dependent on Feature scaled dataset.

## 1.2 Code template in Python

Note that some parts of the code might be omitted.

```

1  # Data Preprocessing Template
2
3  # Importing the libraries
4  import numpy as np
5  import matplotlib.pyplot as plt
6  import pandas as pd
7
8  # Importing the dataset
9  dataset = pd.read_csv('Data.csv')
10 X = dataset.iloc[:, :-1].values
11 y = dataset.iloc[:, 3].values
12
13 # Taking care of missing data
14 from sklearn.preprocessing import Imputer
15 imputer = Imputer(missing_values = 'NaN', strategy = 'mean', axis = 0)
16 imputer = imputer.fit(X[:, 1:3])
17 X[:, 1:3] = imputer.transform(X[:, 1:3])
18
19 # Encoding categorical data
20 # Encoding the Independent Variable
21 from sklearn.preprocessing import LabelEncoder, OneHotEncoder
22 labelencoder_X = LabelEncoder()
23 X[:, 0] = labelencoder_X.fit_transform(X[:, 0])
24 onehotencoder = OneHotEncoder(categorical_features = [0])
25 X = onehotencoder.fit_transform(X).toarray()
26 # Encoding the Dependent Variable
27 labelencoder_y = LabelEncoder()
28 y = labelencoder_y.fit_transform(y)
29
30 # Splitting the dataset into the Training set and Test set
31 from sklearn.cross_validation import train_test_split
32 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)
33
34 # Feature Scaling
35 from sklearn.preprocessing import StandardScaler
36 sc_X = StandardScaler()
37 X_train = sc_X.fit_transform(X_train)
38 X_test = sc_X.transform(X_test)
39 sc_y = StandardScaler()
40 y_train = sc_y.fit_transform(y_train)

```

## 1.3 Code template in R

Note that some parts of the code might be omitted.

```

1  # Data Preprocessing Template
2
3  # Importing the dataset
4  dataset = read.csv('Data.csv')
5
6  # Taking care of missing data
7  dataset$Age = ifelse(is.na(dataset$Age),
8                      ave(dataset$Age, FUN = function(x) mean(x, na.rm = TRUE)),
9                      dataset$Age)
10 dataset$Salary = ifelse(is.na(dataset$Salary),
11                         ave(dataset$Salary, FUN = function(x) mean(x, na.rm = TRUE)),
12                         dataset$Salary)
13
14 # Encoding categorical data
15 dataset$Country = factor(dataset$Country,
16                           levels = c('France', 'Spain', 'Germany'),
17                           labels = c(1, 2, 3))
18 dataset$Purchased = factor(dataset$Purchased,
19                             levels = c('No', 'Yes'),
20                             labels = c(0, 1))
21
22 # Splitting the dataset into the Training set and Test set
23 # install.packages('caTools')
24 library(caTools)
25 set.seed(123)
26 split = sample.split(dataset$DependentVariable, SplitRatio = 0.8)
27 training_set = subset(dataset, split == TRUE)
28 test_set = subset(dataset, split == FALSE)
29
30 # Feature Scaling
31 training_set = scale(training_set)
32 test_set = scale(test_set)

```

## 2 Regression

Regression is a statistical method of predicting the value of **dependent variable (DV)** based on values of **independent variable(s) (IVs)**. More information here [https://en.wikipedia.org/wiki/Regression\\_analysis](https://en.wikipedia.org/wiki/Regression_analysis).

### 2.1 Regression code template in Python

```

1  # Regression
2
3  # Importing the libraries
4  import numpy as np
5  import matplotlib.pyplot as plt
6  import pandas as pd
7
8  # Importing the dataset
9  dataset = pd.read_csv('Data.csv')
10 X = dataset.iloc[:, :-1].values
11 y = dataset.iloc[:, 1].values
12
13 # Splitting the dataset into the Training set and Test set
14 from sklearn.cross_validation import train_test_split
15 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 1/3, random_state = 0)
16
17 # Feature Scaling
18 """from sklearn.preprocessing import StandardScaler
19 sc_X = StandardScaler()
20 X_train = sc_X.fit_transform(X_train)
21 X_test = sc_X.transform(X_test)
22 sc_y = StandardScaler()
23 y_train = sc_y.fit_transform(y_train)"""
24
25 # Fitting Regression to the Training set
26 # Creating Regressor HERE
27
28 # Predicting the Test set results
29 y_pred = regressor.predict(X_test)
30
31 # Visualising the Training set results
32 plt.scatter(X_train, y_train, color = 'red')
33 plt.plot(X_train, regressor.predict(X_train), color = 'blue')
34 plt.title('(Training set)')
35 plt.xlabel('IV')
36 plt.ylabel('DV')
37 plt.show()
38
39 # Visualising the Test set results
40 plt.scatter(X_test, y_test, color = 'red')
41 plt.plot(X_train, regressor.predict(X_train), color = 'blue')
42 plt.title('(Test set)')
43 plt.xlabel('IV')
44 plt.ylabel('DV')
45 plt.show()

```

### 2.2 Regression code template in R

```

1  # Simple Linear Regression
2
3  # Importing the dataset
4  dataset = read.csv('Data.csv')
5
6  # Splitting the dataset into the Training set and Test set
7  # install.packages('caTools')
8  library(caTools)
9  set.seed(123)
10 split = sample.split(dataset$Salary, SplitRatio = 2/3)
11 training_set = subset(dataset, split == TRUE)
12 test_set = subset(dataset, split == FALSE)
13
14 # Feature Scaling
15 # training_set = scale(training_set)
16 # test_set = scale(test_set)
17
18 # Fitting Simple Linear Regression to the Training set
19 # Creating Regressor HERE
20

```

```

21 # Predicting the Test set results
22 y_pred = predict(regressor, newdata = test_set)
23
24 # Visualising the Training set results
25 library(ggplot2)
26 ggplot() +
27   geom_point(aes(x = training_set$YearsExperience, y = training_set$Salary),
28             colour = 'red') +
29   geom_line(aes(x = training_set$YearsExperience, y = predict(regressor, newdata = training_set)),
30            colour = 'blue') +
31   ggtitle('(Training set)') +
32   xlab('IV') +
33   ylab('DV')
34
35 # Visualising the Test set results
36 library(ggplot2)
37 ggplot() +
38   geom_point(aes(x = test_set$YearsExperience, y = test_set$Salary),
39             colour = 'red') +
40   geom_line(aes(x = training_set$YearsExperience, y = predict(regressor, newdata = training_set)),
41            colour = 'blue') +
42   ggtitle('(Test set)') +
43   xlab('IV') +
44   ylab('DV')

```

## 2.3 Simple Linear Regression

**Problem:** Our dataset contains years of experience and salary. We want to fit the model to recommend the salary based on years of experience.

YearsExperience	Salary
1.1	39343.00
1.3	46205.00
1.5	37731.00
2.0	43525.00
2.2	39891.00
2.9	56642.00
3.0	60150.00
3.2	54445.00
3.2	64445.00
3.7	57189.00
3.9	63218.00
4.0	55794.00
4.0	56957.00
4.1	57081.00
4.5	61111.00
4.9	67938.00
5.1	66029.00
5.3	83088.00
5.9	81363.00
6.0	93940.00
6.8	91738.00
7.1	98273.00
7.9	101302.00
8.2	113812.00
8.7	109431.00
9.0	105582.00
9.5	116969.00
9.6	112635.00
10.3	122391.00
10.5	121872.00

Simple Linear Regression is the easiest model following the equation

$$y = b_1 * x + b_0 \quad (1)$$

where  $y$  is the dependent variable which we are interested in (salary) and  $x$  is the independent variable (years of experience).

In order to use Linear Regression we simply replace comment from the template with the following code in Python

```

1 # Fitting Simple Linear Regression to the Training set
2 from sklearn.linear_model import LinearRegression
3 regressor = LinearRegression()
4 regressor.fit(X_train, y_train)

```

And in R we use the following code to fit the Simple Linear Model

```

1 # Fitting Simple Linear Regression to the Training set
2 regressor = lm(formula = Salary ~ YearsExperience,
3               data = training_set)

```

Results are following.



**Fig. 1.** Simple Linear regression in Python

## 2.4 Multiple Linear Regression

**Problem:** We want to predict how profitable will the company be when we know its spending in R&D, Advertisement and Administration. We have a dataset containing 50 startups, their country and their spending to train and test the model. This is the sample of the dataset:

Multiple Linear Regression follows this equation

$$y = b_1 * x_1 + b_2 * x_2 + \dots + b_n * x_n + b_0 \quad (2)$$

where  $y$  is the dependent variable which we are interested in (profit) and  $x_i$  is the independent variable (spending).

### Assumptions of Linear Regression

There are several assumptions of Linear Regression:

- Linearity



- Homoscedasticity  
In statistics, a sequence or a vector of random variables is homoscedastic if all random variables in the sequence or vector have the same finite variance.
- Multivariable normality  
More information: [https://en.wikipedia.org/wiki/Multivariate\\_normal\\_distribution](https://en.wikipedia.org/wiki/Multivariate_normal_distribution)
- Independence of errors  
Error values are statistically independent.
- Lack of multicollinearity  
Multicollinearity in regression occurs when predictor variables (independent variables) in the regression model are more highly correlated with other predictor variables than with the dependent variable. More information: <http://www.researchconsultation.com/multicollinearity-multiple-regression.asp>

## 5 methods of building models

There are some methods to create predicting model. I am listing a few of them.

1. All-in (fit the model with all the independent variable)
2. Backward Elimination
3. Forward Selection
4. Bidirectional Elimination
5. Score comparison (fit the models with all possible combinations of independent variables and compare them)

### Backward elimination

While building a model we might want to get rid of some IVs. Some IVs usually don't have significant impact on the results. Backward elimination is an algorithm that can help us choose only significant variables.

1. Select a significance level to stay in the model (e.g.  $SL = 0.05$ )
2. Fit the full model with all possible predictors
3. Consider the predictor with the highest P-value. If  $P > SL$ , go to step 4, otherwise **Model is ready**
4. Remove the predictor
5. Fit the model without this variable

### Forward Selection

1. Select a significance level to stay in the model (e.g.  $SL = 0.05$ )
2. Fit all simple regression models  $y$   $x_i$  Select the one with the lowest P-value
3. Keep this variable and fit all possible models with one extra predictor added to the one(s) you already have
4. Consider the predictor with the lowest P-value. If  $P < SL$  go to step 3. Otherwise **Your model is ready**

### Bidirectional elimination

1. Select a significance level to enter and to stay in the model (e.g.  $SL_{Enter} = 0.05$ ,  $SL_{Stay} = 0.05$ )
2. Perform the next step of Forward Selection (new variables must have  $P < SL_{Enter}$  to enter)
3. Perform ALL steps of Backward Elimination (old variables must have  $P < SL_{Stay}$  to stay).
4. No new variables can enter and no old variables can exit. **Your model is ready**

### Code

The following code implements Multiple Linear Regression in Python.

```

1 # Fitting Multiple Linear Regression to the Training set
2 from sklearn.linear_model import LinearRegression
3 regressor = LinearRegression()
4 regressor.fit(X_train, y_train)

```

The following code implements Backward elimination in Python. The method called `summary()` is used to display all the P-values of Independent Variables. Then we choose the one with the biggest P-value, remove it and fit the model again.

```

1 # Building the optimal model using Backward Elimination
2 import statsmodels.formula.api as sm
3 X = np.append(arr = np.ones((50, 1)).astype(int), values = X, axis = 1)
4 X_opt = X[:, [0, 1, 2, 3, 4, 5]]
5 regressor_OLS = sm.OLS(endog = y, exog = X_opt).fit()
6 regressor_OLS.summary()
7 X_opt = X[:, [0, 1, 3, 4, 5]]
8 regressor_OLS = sm.OLS(endog = y, exog = X_opt).fit()
9 regressor_OLS.summary()
10 X_opt = X[:, [0, 3, 4, 5]]
11 regressor_OLS = sm.OLS(endog = y, exog = X_opt).fit()
12 regressor_OLS.summary()
13 X_opt = X[:, [0, 3, 5]]
14 regressor_OLS = sm.OLS(endog = y, exog = X_opt).fit()
15 regressor_OLS.summary()
16 X_opt = X[:, [0, 3]]
17 regressor_OLS = sm.OLS(endog = y, exog = X_opt).fit()
18 regressor_OLS.summary()

```

```

"""
                                OLS Regression Results
=====
Dep. Variable:                  y    R-squared:                0.951
Model:                        OLS    Adj. R-squared:           0.945
Method:                    Least Squares    F-statistic:            169.9
Date:                Fri, 02 Jun 2017    Prob (F-statistic):      1.34e-27
Time:                17:27:33    Log-Likelihood:         -525.38
No. Observations:                50    AIC:                   1063.
Df Residuals:                    44    BIC:                   1074.
Df Model:                        5
Covariance Type:                nonrobust
=====
                    coef    std err          t      P>|t|      [0.025     0.975]
-----
const          5.013e+04    6884.820         7.281     0.000     3.62e+04     6.4e+04
x1             198.7888    3371.007         0.059     0.953    -6595.030     6992.607
x2            -41.8870    3256.039        -0.013     0.990    -6604.003     6520.229
x3              0.8060         0.046        17.369     0.000         0.712         0.900
x4             -0.0270         0.052        -0.517     0.608         -0.132         0.078
x5              0.0270         0.017         1.574     0.123         -0.008         0.062
=====
Omnibus:                 14.782    Durbin-Watson:           1.283
Prob(Omnibus):            0.001    Jarque-Bera (JB):        21.266
Skew:                    -0.948    Prob(JB):                2.41e-05
Kurtosis:                 5.572    Cond. No.                1.45e+06
=====
>>>

```

Fig. 2. Calling summary function in Python. We choose the variable with biggest P-value and remove it.

The following code implements Multiple Linear Regression in R.

```

1 # Fitting Multiple Linear Regression to the Training set
2 regressor = lm(formula = Profit ~ .,
3               data = training_set)

```

The following code implements Backward elimination in R.

```

1 # Building the optimal model using Backward Elimination
2 regressor = lm(formula = Profit ~ R.D.Spend + Administration + Marketing.Spend + State,
3               data = dataset)
4 summary(regressor)
5 # Optional Step: Remove State2 only (as opposed to removing State directly)
6 # regressor = lm(formula = Profit ~ R.D.Spend + Administration + Marketing.Spend + factor(State, exclude = 2),
7 #               data = dataset)
8 # summary(regressor)
9 regressor = lm(formula = Profit ~ R.D.Spend + Administration + Marketing.Spend,
10              data = dataset)
11 summary(regressor)
12 regressor = lm(formula = Profit ~ R.D.Spend + Marketing.Spend,
13              data = dataset)
14 summary(regressor)
15 regressor = lm(formula = Profit ~ R.D.Spend,
16              data = dataset)
17 summary(regressor)

```

## 2.5 Polynomial Regression

**Problem:** We have a set of job positions and their according salaries. We want to train the model to propose the best salary for new employees.

Polynomial regression is the first **nonlinear** model.

Position	Level	Salary
Business Analyst	1	45000
Junior Consultant	2	50000
Senior Consultant	3	60000
Manager	4	80000
Country Manager	5	110000
Region Manager	6	150000
Partner	7	200000
Senior Partner	8	300000
C-level	9	500000
CEO	10	1000000

Polynomial Regression follows this equation

$$y = b_1 * x_1^1 + b_2 * x_1^2 + \dots + b_n * x_n^n + b_0 \quad (3)$$

The following code implements the Polynomial regressor in Python

```

1 # Fitting Polynomial Regression to the dataset
2 from sklearn.preprocessing import PolynomialFeatures
3 poly_reg = PolynomialFeatures(degree = 5)
4 X_poly = poly_reg.fit_transform(X)
5 poly_reg.fit(X_poly, y)
6 lin_reg_2 = LinearRegression()
7 lin_reg_2.fit(X_poly, y)

```

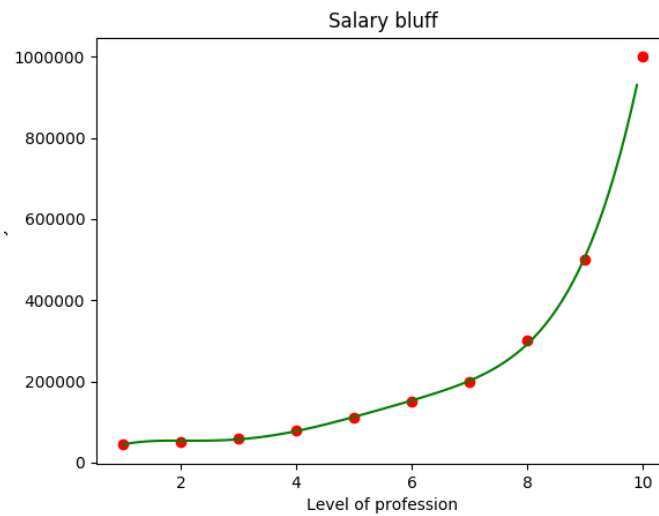
The following code implements the Polynomial regressor in R

```

1 # Fitting Polynomial Regression to the dataset
2 dataset$Level2 = dataset$Level^2
3 dataset$Level3 = dataset$Level^3
4 dataset$Level4 = dataset$Level^4
5 poly_reg = lm(formula = Salary ~ .,
6               data = dataset)

```

The following picture shows visualisation of train model and real data (dots). Model is fitted by polynomial of 5th degree.



**Fig. 3.** Polynomial regression in Python

## 2.6 Support Vector Regression (SVR)

**Problem:** The same as in Polynomial regression.

SVR is the second **nonlinear** model. SVR model from `sklearn` Python library requires the dataset to be scaled in preprocessing.

The following code implements Support Vector Regressor in Python

```

1 # Fitting SVR to the dataset
2 from sklearn.svm import SVR
3 regressor = SVR(kernel = 'rbf')
4 regressor.fit(X, y)

```

*Kernel* is the method used by the predictor. *RBF* means *Radial Basis function kernel*. More information here: [https://en.wikipedia.org/wiki/Kernel\\_method](https://en.wikipedia.org/wiki/Kernel_method)

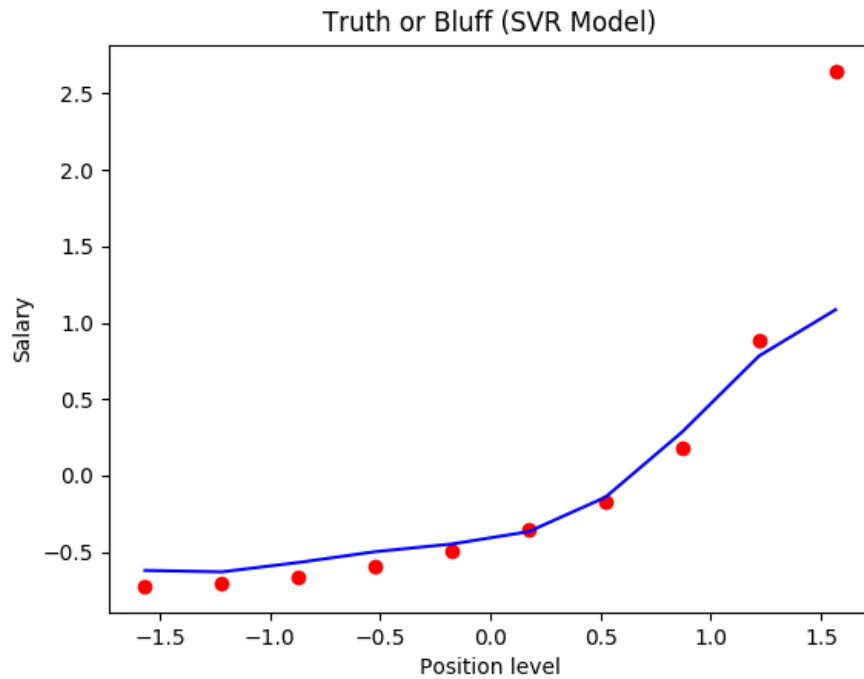


Fig. 4. Support Vector regression in Python

## 2.7 Decision Tree Regression

**Problem:** The same as in Polynomial regression.

The following code implements Decision Tree Regressor in Python

```

1 # Fitting Decision Tree Regression to the dataset
2 from sklearn.tree import DecisionTreeRegressor
3 regressor = DecisionTreeRegressor(random_state = 0)
4 regressor.fit(X, y)

```

The following code implements Decision Tree Regressor in R

```

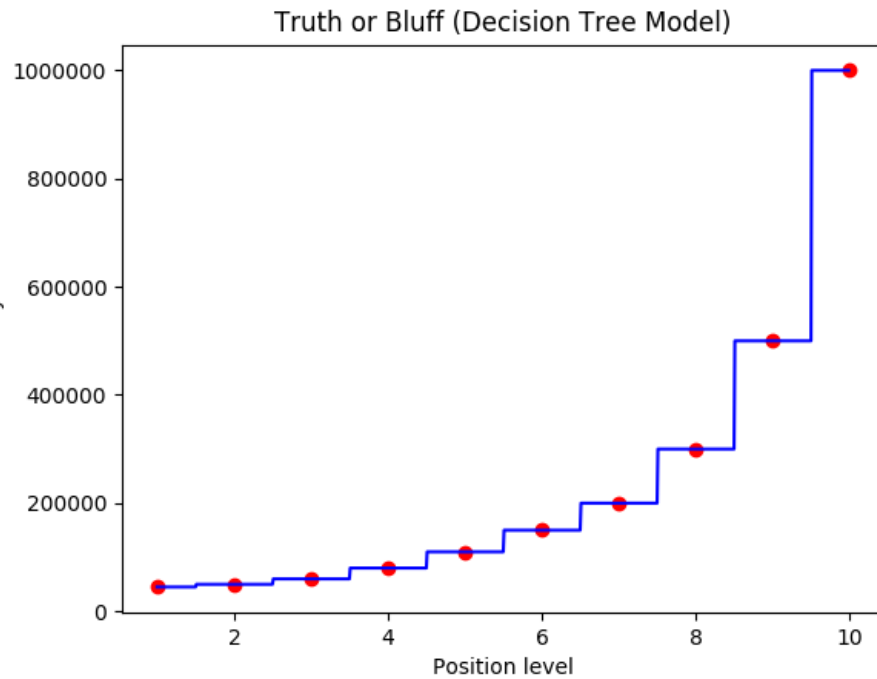
1 # Fitting Decision Tree Regression to the dataset
2 # install.packages('rpart')
3 library(rpart)

```

```

4 regressor = rpart(formula = Salary ~ .,
5                   data = dataset,
6                   control = rpart.control(minsplit = 1))

```



**Fig. 5.** Decision Tree regression in Python

## 2.8 Random Forest Regression

**Problem:** The same as in Polynomial regression.

Random Forest algorithm uses multiple Decision Trees and then average the results. Number of Decision Trees determines the number of splits and the value level of prediction.

The following code implements Random Forest Regressor in Python

```

1 # Fitting Random Forest Regression to the dataset
2 from sklearn.ensemble import RandomForestRegressor

```

```

3 regressor = RandomForestRegressor(n_estimators = 10, random_state
4 regressor.fit(X, y)

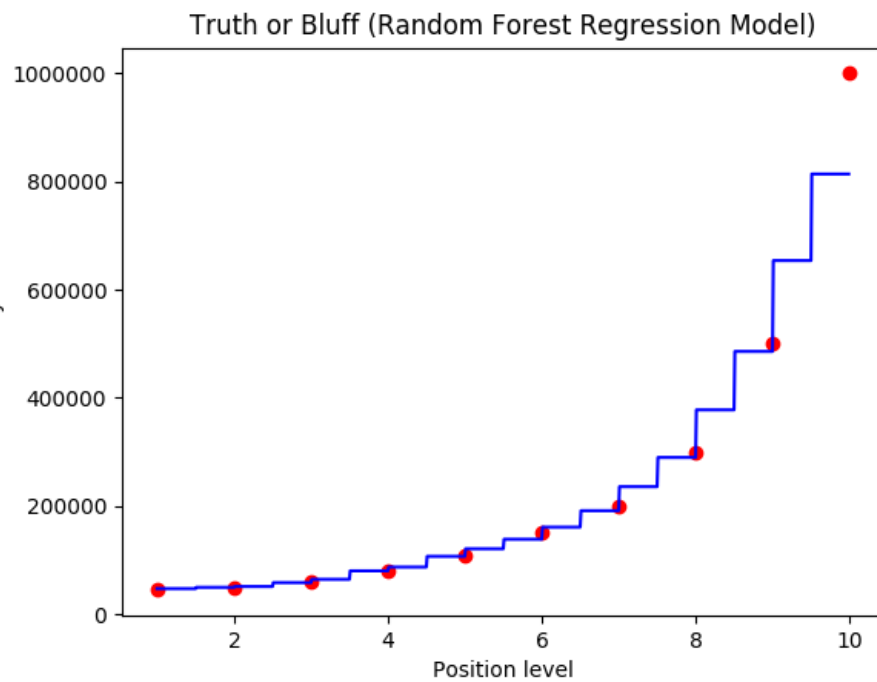
```

The following code implements Random Forest Regressor in R

```

1 # Fitting Random Forest Regression to the dataset
2 # install.packages('randomForest')
3 library(randomForest)
4 set.seed(1234)
5 regressor = randomForest(x = dataset[-2],
6                           y = dataset$Salary,
7                           ntree = 500)

```



**Fig. 6.** Random Forest regression in Python



## 2.9 Evaluating model performance

### R-squared

$R^2$  - The Goodness of fit

$$R^2 = 1 - \frac{SS_{res}}{SS_{tot}} \quad (4)$$

$$SS_{res} = \sum_0^n (y_{i_{actual}} - y_{i_{predicted}})^2$$

$$SS_{tot} = \sum_0^n (y_{i_{actual}} - y_{average})^2$$

$SS_{res}$  means sum of squares residual and  $SS_{tot}$  means sum of squares total,  $n$  is the number of observations.

The closer  $R^2$  is to 1, the better.

The biggest disadvantage is that  $R^2$  never decrease if we introduce new independent variable to the model. That's why we use Adjusted  $R^2$ .

### Adjusted R-squared

$$R_{adj}^2 = 1 - (1 - R^2) \cdot \frac{n - 1}{n - p - 1} \quad (5)$$

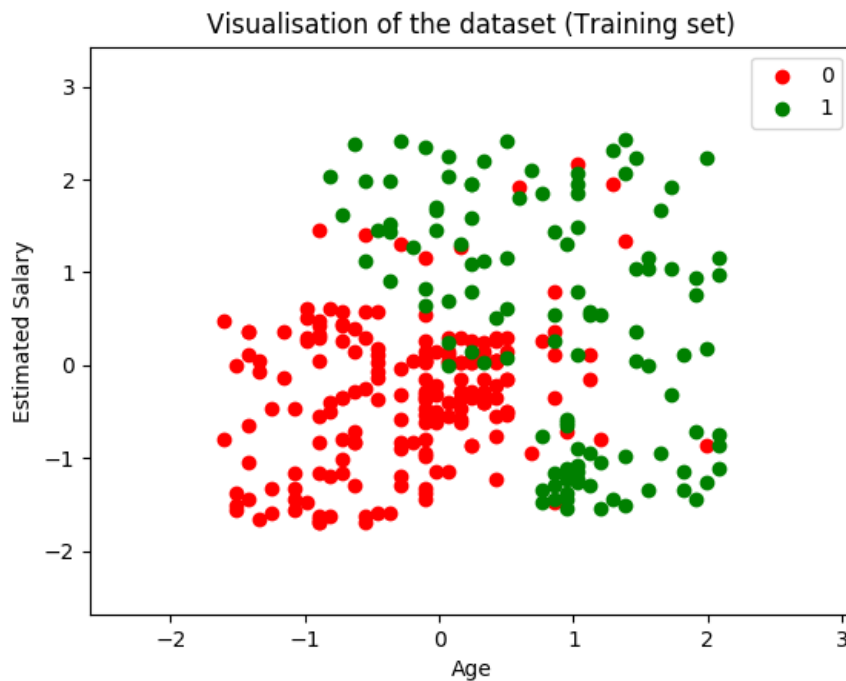
$p$ ... number of regressors (independent variables),  $n$ ... sample size

### 3 Classification

Classification is a statistical method which identifies a new observation into one of the several categories. More information here: [https://en.wikipedia.org/wiki/Statistical\\_classification](https://en.wikipedia.org/wiki/Statistical_classification).

**Problem:** We will try to deal with the same problem for all listed methods below. We have a dataset of 400 users of some social network. The dataset tells us whether the user bought or didn't buy the product. The dataset contains user ID, salary and age. This is the sample of the dataset and the visualisation of the training set. Note that the axis shows values already scaled

User ID	Gender	Age	EstimatedSalary	Purchased
15624510	Male	19	19000	0
15810944	Male	35	20000	0
15668575	Female	26	43000	0
15569641	Female	58	95000	1
15815236	Female	45	131000	1
15811177	Female	35	77000	0
15680587	Male	36	144000	1



**Fig. 7.** Visualisation of the dataset for Classification problems. 1 means purchased, 0 means otherwise. Python

### 3.1 Classification code template in Python

```

1  # Classification template
2
3  # Importing the libraries
4  import numpy as np
5  import matplotlib.pyplot as plt
6  import pandas as pd
7
8  # Importing the dataset
9  dataset = pd.read_csv('Social_Network_Ads.csv')
10 X = dataset.iloc[:, [2, 3]].values
11 y = dataset.iloc[:, 4].values
12
13 # Splitting the dataset into the Training set and Test set
14 from sklearn.cross_validation import train_test_split
15 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state = 0)
16
17 # Feature Scaling
18 from sklearn.preprocessing import StandardScaler
19 sc = StandardScaler()
20 X_train = sc.fit_transform(X_train)
21 X_test = sc.transform(X_test)
22
23 # Fitting classifier to the Training set
24 # Create your classifier here
25
26 # Predicting the Test set results

```

```

27 y_pred = classifier.predict(X_test)
28
29 # Making the Confusion Matrix
30 from sklearn.metrics import confusion_matrix
31 cm = confusion_matrix(y_test, y_pred)
32
33 # Visualising the Training set results
34 from matplotlib.colors import ListedColormap
35 X_set, y_set = X_train, y_train
36 X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[:, 0].max() + 1, step = 0.01),
37                      np.arange(start = X_set[:, 1].min() - 1, stop = X_set[:, 1].max() + 1, step = 0.01))
38 plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()]).T).reshape(X1.shape),
39             alpha = 0.75, cmap = ListedColormap(('red', 'green')))
40 plt.xlim(X1.min(), X1.max())
41 plt.ylim(X2.min(), X2.max())
42 for i, j in enumerate(np.unique(y_set)):
43     plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
44               c = ListedColormap(('red', 'green'))(i), label = j)
45 plt.title('Classifier (Training set)')
46 plt.xlabel('Age')
47 plt.ylabel('Estimated Salary')
48 plt.legend()
49 plt.show()
50
51 # Visualising the Test set results
52 from matplotlib.colors import ListedColormap
53 X_set, y_set = X_test, y_test
54 X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[:, 0].max() + 1, step = 0.01),
55                      np.arange(start = X_set[:, 1].min() - 1, stop = X_set[:, 1].max() + 1, step = 0.01))
56 plt.contourf(X1, X2, classifier.predict(np.array([X1.ravel(), X2.ravel()]).T).reshape(X1.shape),
57             alpha = 0.75, cmap = ListedColormap(('red', 'green')))
58 plt.xlim(X1.min(), X1.max())
59 plt.ylim(X2.min(), X2.max())
60 for i, j in enumerate(np.unique(y_set)):
61     plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],
62               c = ListedColormap(('red', 'green'))(i), label = j)
63 plt.title('Classifier (Test set)')
64 plt.xlabel('Age')
65 plt.ylabel('Estimated Salary')
66 plt.legend()
67 plt.show()

```

### 3.2 Classification code template in R

```

1 # Classification template
2
3 # Importing the dataset
4 dataset = read.csv('Social_Network_Ads.csv')
5 dataset = dataset[3:5]
6
7 # Encoding the target feature as factor
8 dataset$Purchased = factor(dataset$Purchased, levels = c(0, 1))
9
10 # Splitting the dataset into the Training set and Test set
11 # install.packages('caTools')
12 library(caTools)
13 set.seed(123)
14 split = sample.split(dataset$Purchased, SplitRatio = 0.75)
15 training_set = subset(dataset, split == TRUE)
16 test_set = subset(dataset, split == FALSE)
17
18 # Feature Scaling
19 training_set[-3] = scale(training_set[-3])
20 test_set[-3] = scale(test_set[-3])
21
22 # Fitting classifier to the Training set
23 # Create your classifier here
24
25 # Predicting the Test set results
26 y_pred = predict(classifier, newdata = test_set[-3])
27
28 # Making the Confusion Matrix
29 cm = table(test_set[, 3], y_pred)
30
31 # Visualising the Training set results
32 library(ElemStatLearn)
33 set = training_set
34 X1 = seq(min(set[, 1]) - 1, max(set[, 1]) + 1, by = 0.01)
35 X2 = seq(min(set[, 2]) - 1, max(set[, 2]) + 1, by = 0.01)
36 grid_set = expand.grid(X1, X2)
37 colnames(grid_set) = c('Age', 'EstimatedSalary')
38 y_pred = predict(classifier, newdata = grid_set)
39 plot(set[, -3],
40      main = 'Classifier (Training set)',
41      xlab = 'Age', ylab = 'Estimated Salary',
42      xlim = range(X1), ylim = range(X2))

```

```

43 contour(X1, X2, matrix(as.numeric(y_grid), length(X1), length(X2)), add = TRUE)
44 points(grid_set, pch = '.', col = ifelse(y_grid == 1, 'springgreen3', 'tomato'))
45 points(set, pch = 21, bg = ifelse(set[, 3] == 1, 'green4', 'red3'))
46
47 # Visualising the Test set results
48 library(ElemStatLearn)
49 set = test_set
50 X1 = seq(min(set[, 1]) - 1, max(set[, 1]) + 1, by = 0.01)
51 X2 = seq(min(set[, 2]) - 1, max(set[, 2]) + 1, by = 0.01)
52 grid_set = expand.grid(X1, X2)
53 colnames(grid_set) = c('Age', 'EstimatedSalary')
54 y_grid = predict(classifier, newdata = grid_set)
55 plot(set[, -3], main = 'Classifier (Test set)',
56      xlab = 'Age', ylab = 'Estimated Salary',
57      xlim = range(X1), ylim = range(X2))
58 contour(X1, X2, matrix(as.numeric(y_grid), length(X1), length(X2)), add = TRUE)
59 points(grid_set, pch = '.', col = ifelse(y_grid == 1, 'springgreen3', 'tomato'))
60 points(set, pch = 21, bg = ifelse(set[, 3] == 1, 'green4', 'red3'))

```

### 3.3 Logistic regression

Logistic regression deals with the situations when DV is in binary form (0/1, bought/didn't bought, etc.). The basic theory of Logistic regression is as follows: We take the Simple Linear Regression

$$y = b_0 + b_1 * x_1, \quad (6)$$

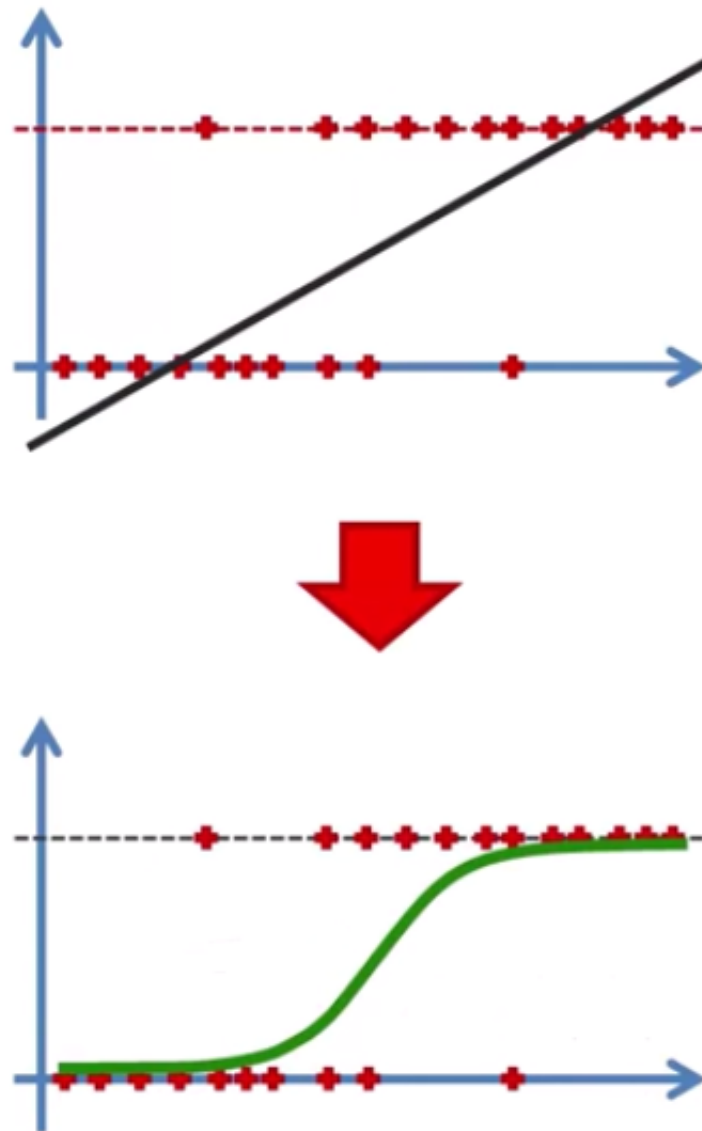
apply the *Sigmoid function*

$$p = \frac{1}{1 + e^{-y}} \quad (7)$$

and solve

$$\ln\left(\frac{p}{1-p}\right) = b_0 + b_1 * x_1. \quad (8)$$

So basically this is what happens:



**Fig. 8.** Top - Simple Linear Regression, Bottom - Logistic Regression

This code implements Logistic regression in Python.

```

1 # Fitting Logistic Regression to the Training set
2 from sklearn.linear_model import LogisticRegression
3 classifier = LogisticRegression(random_state = 0)
4 classifier.fit(X_train, y_train)

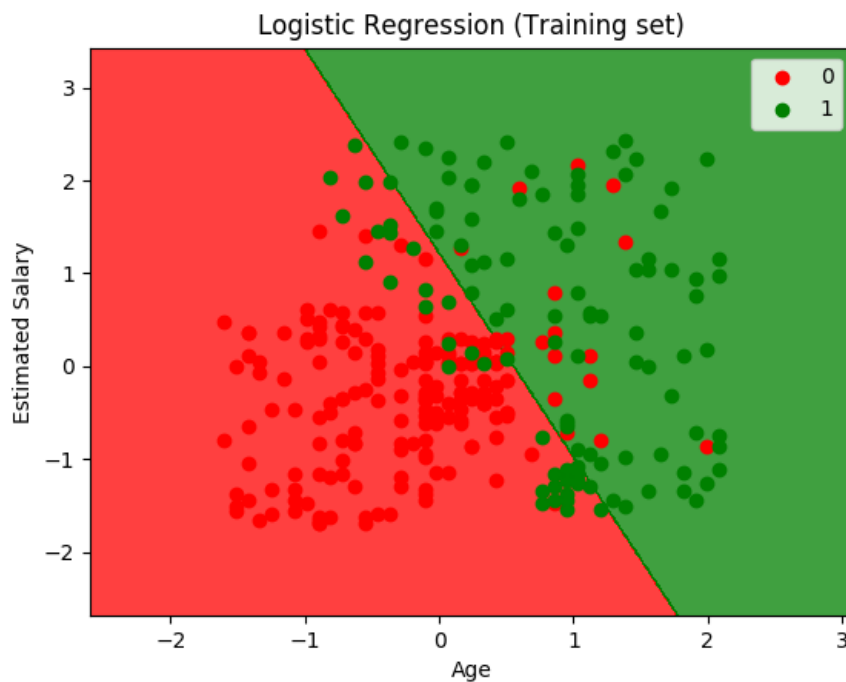
```

This code implements Logistic regression in R.

```

1 # Fitting Logistic Regression to the Training set
2 classifier = glm(formula = Purchased ~ .,
3                  family = binomial,
4                  data = training_set)

```



**Fig. 9.** Logistic regression in Python

### 3.4 K-Nearest neighbours

K-Nearest neighbours works as follows: we take the new observation which we want to classify. Then we take a look at  $k$  nearest points from the training set. Based on the nearest neighbours we decide the class of the new observation. We need to choose  $k$  first.

This code implements the K-NN classifier in Python.

```

1 # Fitting K-NN to the Training set
2 from sklearn.neighbors import KNeighborsClassifier
3 classifier = KNeighborsClassifier(n_neighbors = 5,
4     metric = 'minkowski', p = 2)
5 classifier.fit(X_train, y_train)

```

This code implements the K-NN classifier in R.

```

1 # Fitting K-NN to the Training set and Predicting the Test set res
2 library(class)
3 y_pred = knn(train = training_set[, -3],
4     test = test_set[, -3],
5     cl = training_set[, 3],
6     k = 5,
7     prob = TRUE)

```



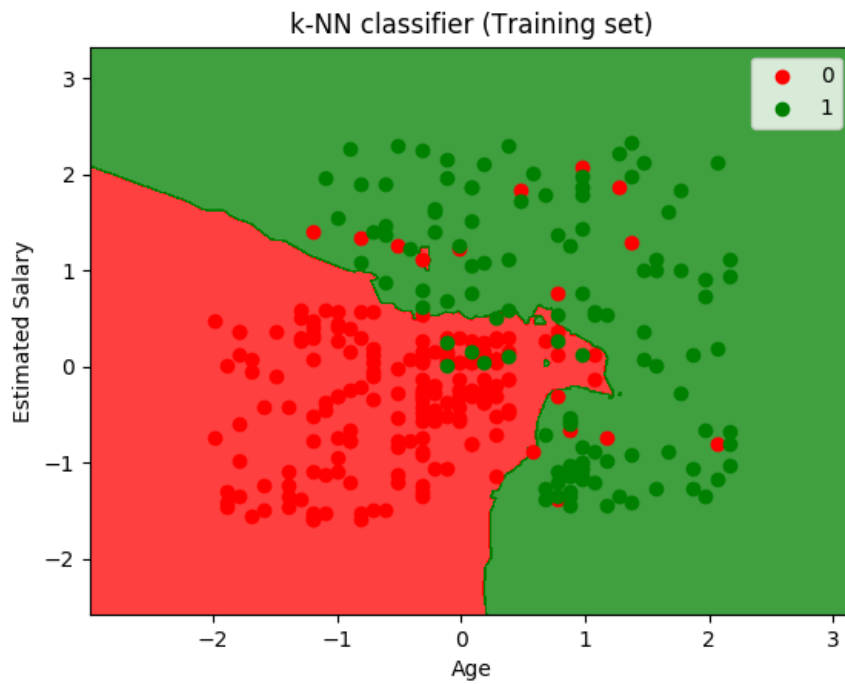
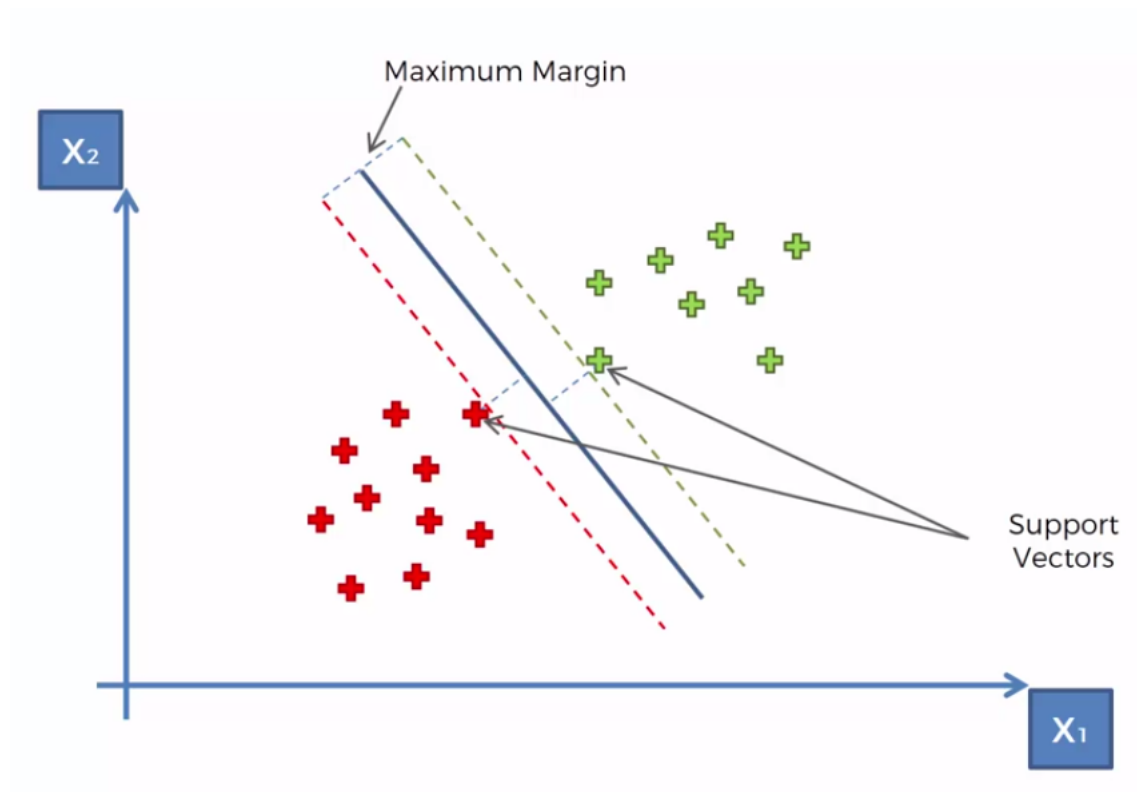


Fig. 10. k-NN algorithm with  $k = 5$  in Python

### 3.5 SVM (Support Vector Machine)

SVM classifier can be linear with linear kernel. SVM is **nonlinear** when using different kernel using the *kernel trick* which maps the data to higher dimension to be linearly separable. More about SVM: [https://en.wikipedia.org/wiki/Support\\_vector\\_machine](https://en.wikipedia.org/wiki/Support_vector_machine).



**Fig. 11.** Support Vector Machine basic

This code implements the SVM classifier in Python.

```

1 # Fitting SVM to the Training set
2 from sklearn.svm import SVC
3 classifier = SVC(kernel = 'linear', random_state = 0)
4 classifier.fit(X_train, y_train)

```

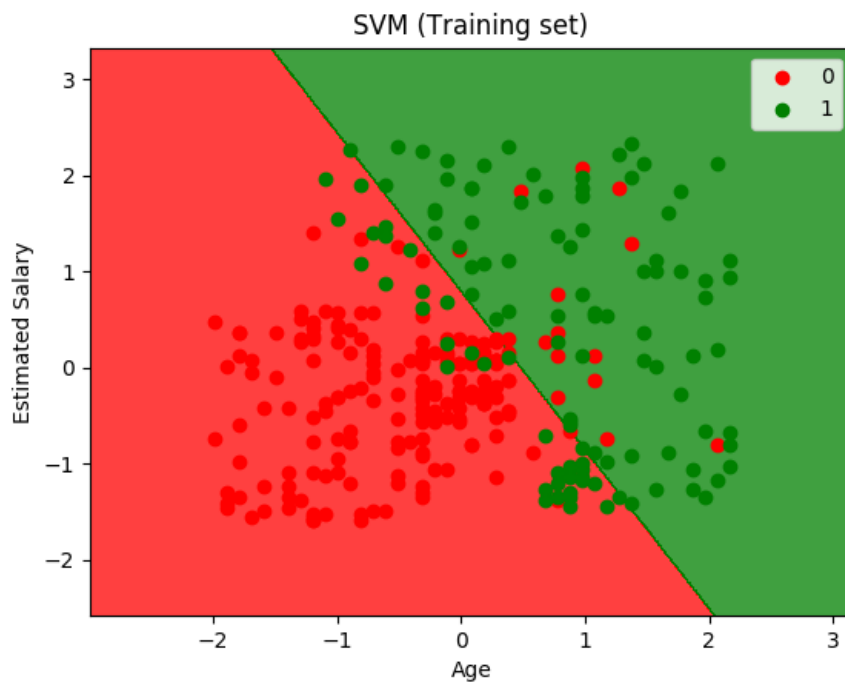
This code implements the SVM classifier in R.

```

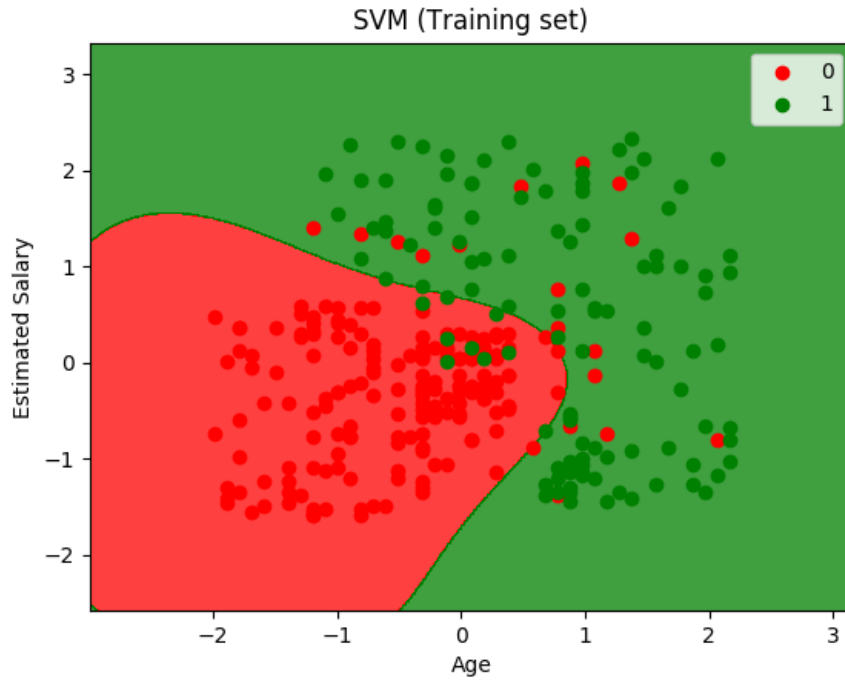
1 # Fitting SVM to the Training set
2 # install.packages('e1071')
3 library(e1071)
4 classifier = svm(formula = Purchased ~ .,

```

```
5 data = training_set ,  
6 type = 'C-classification',  
7 kernel = 'linear')
```



**Fig. 12.** Support Vector Machine with **linear** kernel in Python



**Fig. 13.** Support Vector Machine with **RBF** kernel in Python

### 3.6 Kernel SVM

More about Kernel SVM in the Udemy course.

### 3.7 Naive Bayes

The Bayes theorem is the known equation:

$$P(A|X) = \frac{P(X|A) \cdot P(A)}{P(X)}. \quad (9)$$

$A$  is the DV and can have two values  $a_1; a_2$  and  $X$  is the feature vector.

In Naive Bayes Classification method we calculate the variables in the equation in this order:

1.  $P(A) = \frac{\text{number of observations with property-}a_i}{\text{number of total observations}}$  ... prior probability.
2.  $P(X) = \frac{\text{number of similar observations}}{\text{number of total observations}}$  ... marginal likelihood. Probability of the occurrence of the particular combination of feature vector  $X$ . Number of similar observations depends on the area we define.

3.  $P(X|A) = \frac{\text{number of similar observations with property-}a_i}{\text{number of total observations with property-}a_i} \dots \text{likelihood.}$
4.  $P(A|X)$  ... finally we get posterior probability which we were looking for.

This code implements the Naive Bayes Classifier in Python.

```

1 # Fitting Naive Bayes to the Training set
2 from sklearn.naive_bayes import GaussianNB
3 classifier = GaussianNB()
4 classifier.fit(X_train, y_train)

```

This code implements the Naive Bayes Classifier in R.

```

1 # Fitting Naive Bayes to the Training set
2 # install.packages('e1071')
3 library(e1071)
4 classifier = naiveBayes(x = training_set[-3],
5                         y = training_set$Purchased)

```

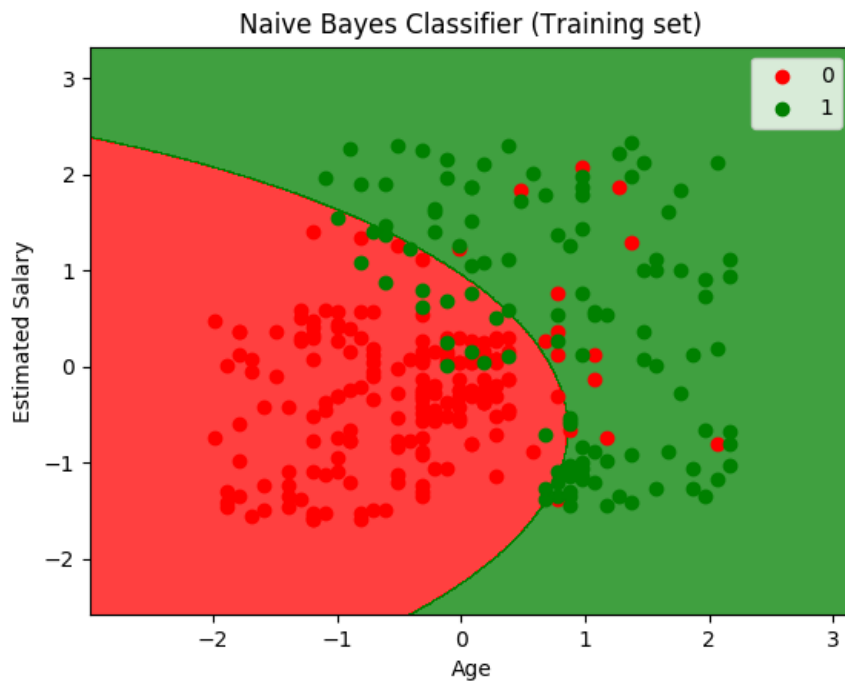


Fig. 14. Naive Bayes Classifier in Python

### 3.8 Decision Tree Classification

Term **CART** - Classification and Regression Trees.

Similar to Decision Tree Regression.

This code implements the Decision Tree Classifier in Python.

```

1 # Fitting Decision Tree Classification to the Training set
2 from sklearn.tree import DecisionTreeClassifier
3 classifier = DecisionTreeClassifier(criterion = 'entropy')
4 classifier.fit(X_train, y_train)

```

This code implements the Decision Tree Classifier in R.

```

1 # Fitting Decision Tree Classification to the Training set
2 # install.packages('rpart')
3 library(rpart)

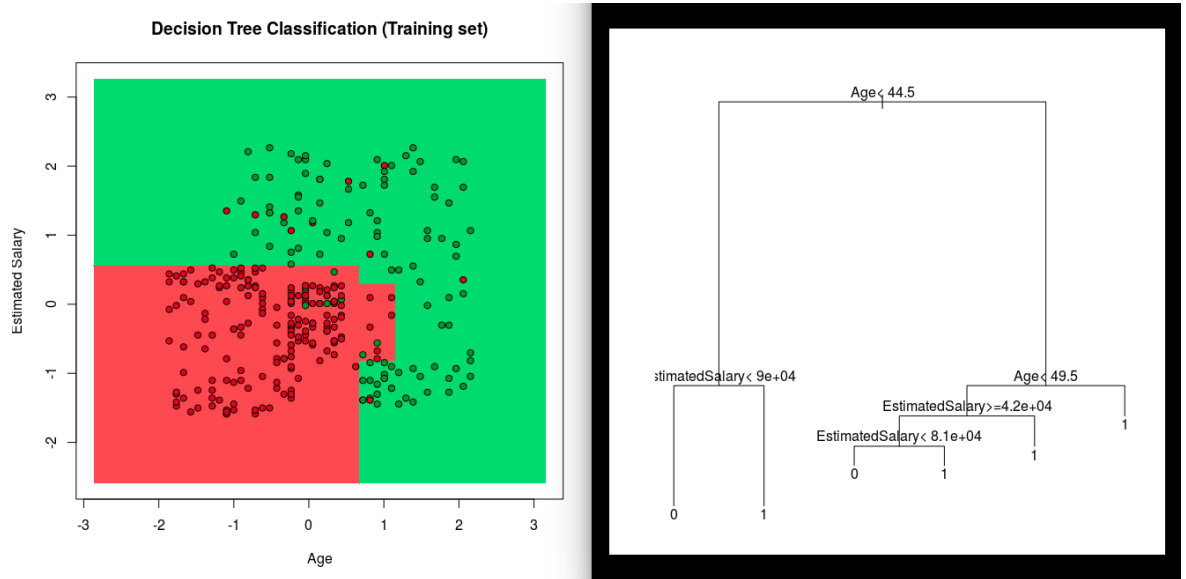
```

```
4 classifier = rpart(formula = Purchased ~ .,  
5                   data = training_set)
```



**Fig. 15.** Decision Tree in Python

In R we can also see the Decision Tree itself simply by using `plot(classifier)`.



**Fig. 16.** Decision Tree in R. Visualisation and Decision Tree itself

### 3.9 Random Forest Classification

Similarly to Random Forest Regression, the Random Forest Classification uses many Decision Trees and then average their results. Random Forest Classification is the algorithm of choice for the Microsoft Kinect developers.

This code implements the Decision Tree Classifier in Python.

```

1 # Fitting Random Forest Classification to the Training set
2 from sklearn.ensemble import RandomForestClassifier
3 classifier = RandomForestClassifier(n_estimators = 10,
4                                   criterion = 'entropy')
5 classifier.fit(X_train, y_train)

```

This code implements the Decision Tree Classifier in R.

```

1 # Fitting Random Forest Classification to the Training set
2 # install.packages('randomForest')
3 library(randomForest)
4 set.seed(123)
5 classifier = randomForest(x = training_set[-3],

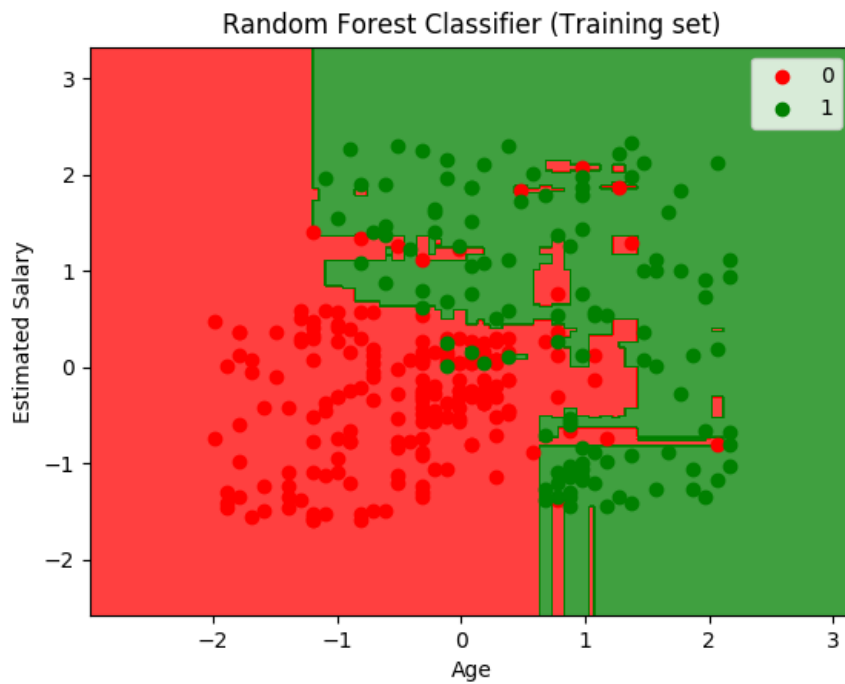
```



```

6 y = training_set$Purchased ,
7 ntree = 500)

```



**Fig. 17.** Random Forest Classifier with 10 Decision Trees in Python

### 3.10 Evaluating Classification model performance

The following approaches are used to pick the best model.

#### Confusion matrix

Confusion matrix tells us number of Correct predictions (sum of diagonal) and number of wrong predictions. It also distinguish

- Type error 1 (False positive)
- Type error 2 (False negative).

		Condition (as determined by "Gold standard")			
		Condition Positive	Condition Negative		
Test Outcome	Test Outcome Positive	True Positive	False Positive (Type I error)	Positive predictive value = $\frac{\Sigma \text{ True Positive}}{\Sigma \text{ Test Outcome Positive}}$	
	Test Outcome Negative	False Negative (Type II error)	True Negative	Negative predictive value = $\frac{\Sigma \text{ True Negative}}{\Sigma \text{ Test Outcome Negative}}$	
		Sensitivity = $\frac{\Sigma \text{ True Positive}}{\Sigma \text{ Condition Positive}}$	Specificity = $\frac{\Sigma \text{ True Negative}}{\Sigma \text{ Condition Negative}}$		

Fig. 18. Confusion matrix

### Accuracy paradox

Unfortunately the confusion matrix has one big withdrawal. Let's define Accuracy rate.

$$AR = \frac{n_{correct}}{n_{total}} \quad (10)$$

It should not be considered the ultimate measure for the performance of the models. Let's consider the following scenario. We build a model with the following confusion matrix:

Table 1. Situation before

	Predicted DV	
	0	1
Actual DV	9700	150
	50	100

We simply calculate that the Accuracy rate is  $AR = 9800/10000 = 98\%$ . Then we abandon the model and always predict negative (0). We come out with the following confusion matrix.

**Table 2.** Situation without model but with better AR

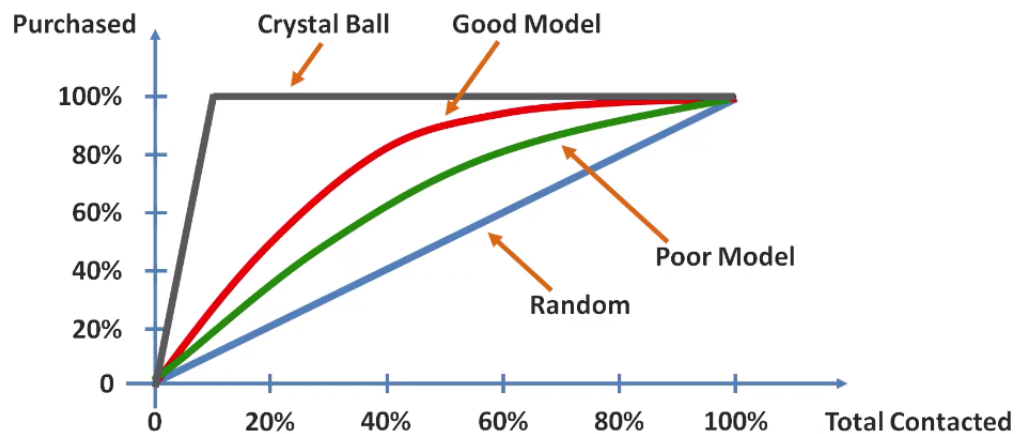
	Predicted DV	
	0	1
	0	1
Actual DV	0 9850	0
	1 150	0

We calculate that the Accuracy rate is  $AR = 9850/10000 = 98.5\%$ . We have three times more Type 2 errors but the AR is better. That's why **Accuracy rate is not the best evaluating technique!**

### CAP curve

**CAP** stands for Cumulative Accuracy Profile.

**Problem:** We have a dataset of 100000 customers and 10% of them will purchase the product. We want to build a model which finds the customers with the highest probability of purchasing the product. Then we contact them. We want our model to have the highest possible hit rate of contacted-purchased. The "Crystal ball" will have 100% hit rate, that means first 10000 contacted equals 10000 purchases.

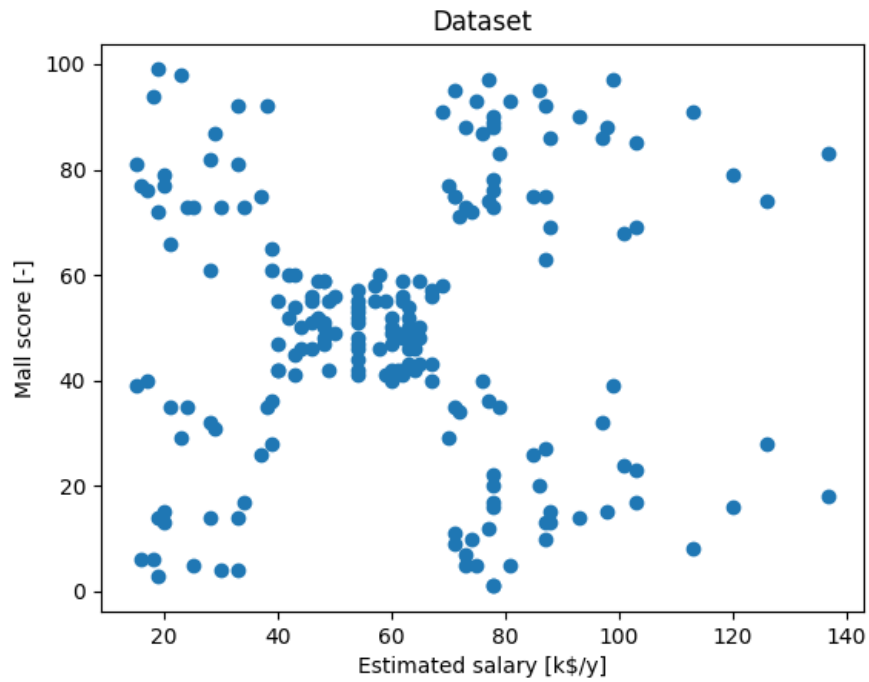
**Fig. 19.** CAP curve explained

## 4 Clustering

Clustering is the form of unsupervised learning. That means we do not have a training set but we have to find the clusters ourselves. We will show two approaches. The first step in both of them is to establish the number of clusters  $k$ . Each approach use different method.

**Problem:** We are data analysts hired by mall company. They provide us with the dataset containing information about their customers. They want us to divide their customers into several (not specified  $k$ ) classes. This is the sample of the dataset. We will use only *salary* and *spending score*.

CustomerID	Genre	Age	Annual Income (k)	Spending Score (1-100)
0001	Male	19	15	39
0002	Male	21	15	81
0003	Female	20	16	6
0004	Female	23	16	77
0005	Female	31	17	40
0006	Female	22	17	76



**Fig. 20.** Visualisation of the dataset for clustering problem in Python

#### 4.1 K-means clustering

1. Choose the number  $k$  of clusters.
2. Select at random  $k$  points. The *centroids*. (not necessarily from the dataset)
3. Assign each datapoint to the closest *centroid*, that forms  $k$  clusters.
4. Compute and place the new *centroids* for each cluster.
5. Reassigning each point according to the new *centroids* and go to step 4. If no reassignment took place. **Your model is ready.**

Selecting the random points for *centroids* might be problematic and can lead to faulty results. K-means++ algorithm deals with it.

#### Elbow method to find optimal K

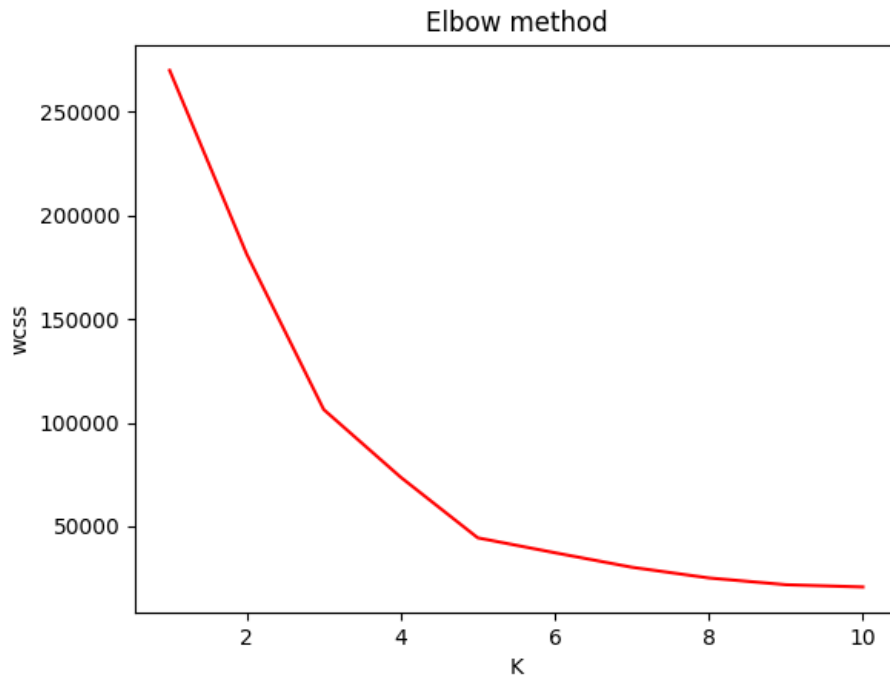
First of all we need to establish the number of clusters  $k$ . In K-means we use the "Elbow method".

We place a *Centroid* in the dataset and calculate the distances between all the points and the centroid. Then we place another *centroid* and calculate and sum the distances of all points to the closest *centroids*. We will call this value *WCSS*. We continue to increase the number of *centroids*  $k$ . It is clear that as  $k$  closes to the number of the points in the dataset, the *WCSS* gets closer to zero.

The equation can be written as

$$WCSS = \sum_{i=1}^k \left( \sum_{P_i \text{ in Cluster } i} dist(P_i, C_i)^2 \right) \quad (11)$$

where  $k$  is the number of clusters. We can plot the values in the graph and choose the  $k$  from the "elbow".



**Fig. 21.** Elbow method to determine the optimal K in Python

From the picture we see that the optimal  $k$  is 5.

### Implementation

This code implements the K-means algorithm in Python with the elbow method and visualisation.

```

1  # K-Means Clustering
2
3  # Importing the libraries
4  import numpy as np
5  import matplotlib.pyplot as plt
6  import pandas as pd
7
8  # Importing the dataset
9  dataset = pd.read_csv('Mall_Customers.csv')
10 X = dataset.iloc[:, [3, 4]].values
11
12 # Using the elbow method to find the optimal number of clusters
13 from sklearn.cluster import KMeans
14 wcss = []
15 for i in range(1, 11):
16     kmeans = KMeans(n_clusters = i, init = 'k-means++', random_state = 42)
17     kmeans.fit(X)
18     wcss.append(kmeans.inertia_)
19 plt.plot(range(1, 11), wcss)

```

```

20 plt.title('The Elbow Method')
21 plt.xlabel('Number of clusters')
22 plt.ylabel('WCSS')
23 plt.show()
24
25 # Fitting K-Means to the dataset
26 kmeans = KMeans(n_clusters = 5, init = 'k-means++', random_state = 42)
27 y_kmeans = kmeans.fit_predict(X)
28
29 # Visualising the clusters
30 plt.scatter(X[y_kmeans == 0, 0], X[y_kmeans == 0, 1], s = 100, c = 'red', label = 'Cluster 1')
31 plt.scatter(X[y_kmeans == 1, 0], X[y_kmeans == 1, 1], s = 100, c = 'blue', label = 'Cluster 2')
32 plt.scatter(X[y_kmeans == 2, 0], X[y_kmeans == 2, 1], s = 100, c = 'green', label = 'Cluster 3')
33 plt.scatter(X[y_kmeans == 3, 0], X[y_kmeans == 3, 1], s = 100, c = 'cyan', label = 'Cluster 4')
34 plt.scatter(X[y_kmeans == 4, 0], X[y_kmeans == 4, 1], s = 100, c = 'magenta', label = 'Cluster 5')
35 plt.scatter(kmeans.cluster_centers_[0, 0], kmeans.cluster_centers_[0, 1],
36             s = 300, c = 'yellow', label = 'Centroids')
37 plt.title('Dataset clustered by HC')
38 plt.ylabel('Mall score [-]')
39 plt.xlabel('Estimated salary [k$/y]')
40 plt.legend()
41 plt.show()

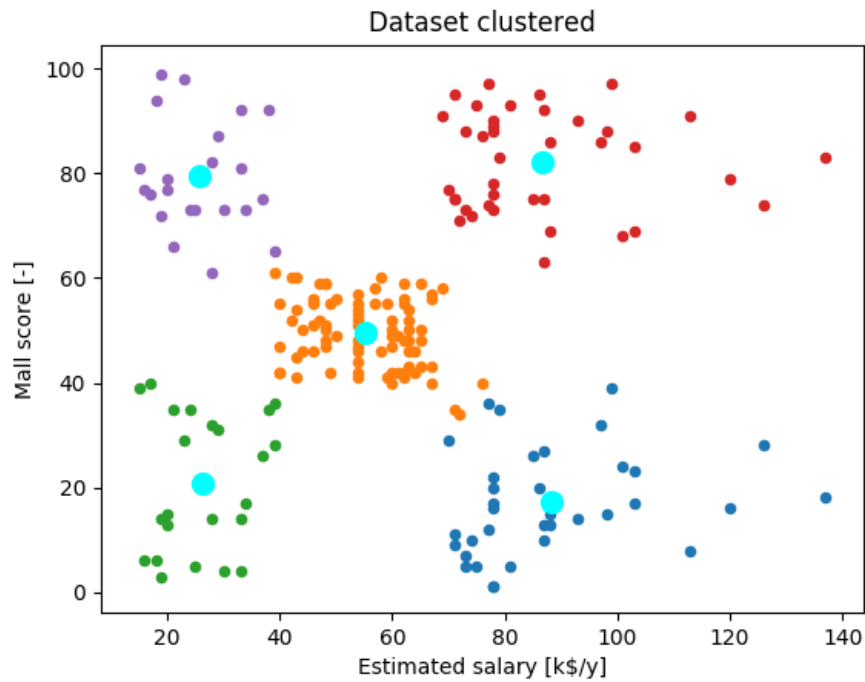
```

This code implements the K-means algorithm in R with the elbow method and visualisation.

```

1 # K-Means Clustering
2
3 # Importing the dataset
4 dataset = read.csv('Mall_Customers.csv')
5 dataset = dataset[4:5]
6
7 # Using the elbow method to find the optimal number of clusters
8 set.seed(6)
9 wcss = vector()
10 for (i in 1:10) wcss[i] = sum(kmeans(dataset, i)$withinss)
11 plot(1:10,
12      wcss,
13      type = 'b',
14      main = paste('The Elbow Method'),
15      xlab = 'Number of clusters',
16      ylab = 'WCSS')
17
18 # Fitting K-Means to the dataset
19 set.seed(29)
20 kmeans = kmeans(x = dataset, centers = 5)
21 y_kmeans = kmeans$cluster
22
23 # Visualising the clusters
24 library(cluster)
25 clusplot(dataset,
26          y_kmeans,
27          lines = 0,
28          shade = TRUE,
29          color = TRUE,
30          labels = 2,
31          plotchar = FALSE,
32          span = TRUE,
33          main = paste('Clusters of customers'),
34          xlab = 'Annual Income',
35          ylab = 'Spending Score')

```



**Fig. 22.** Result of K-mean algorithm in Python, blue points are centroids.

## 4.2 Hierarchical clustering

Agglomerative HC:

1. Make each data point a single cluster (that makes  $N$  clusters)
2. Take the two closest points and make them 1 cluster ( $N-1$  clusters)
3. Take the two **closest clusters** and make them 1 cluster ( $N-2$  clusters)
4. Repeat step 3 until there is only 1 cluster

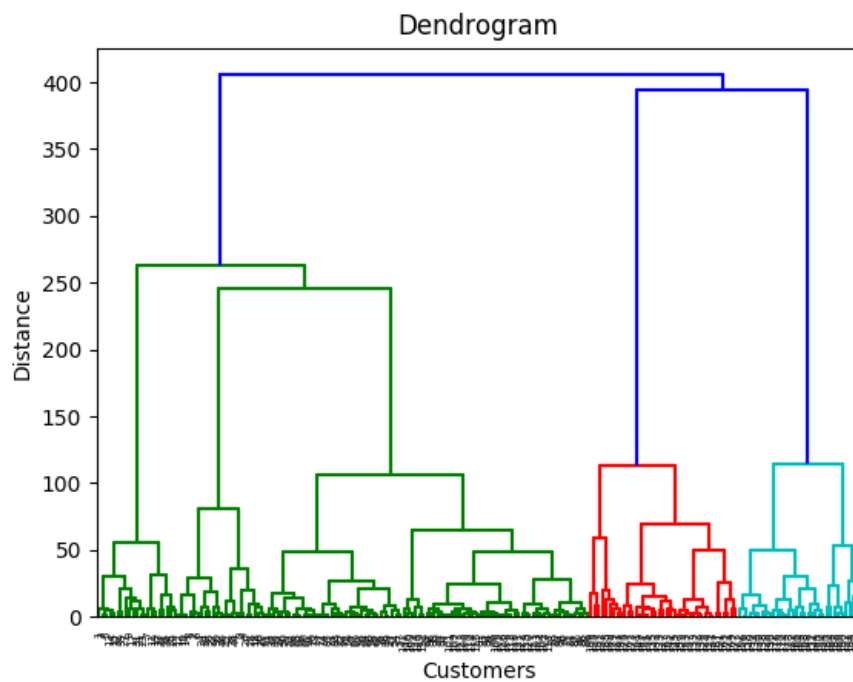
What is the **distance between clusters**? We have several valid answers. It depends on the usecase.

- Distance of closest points
- Distance of furthest points
- Average distance
- Distance of centroids



## Dendrograms

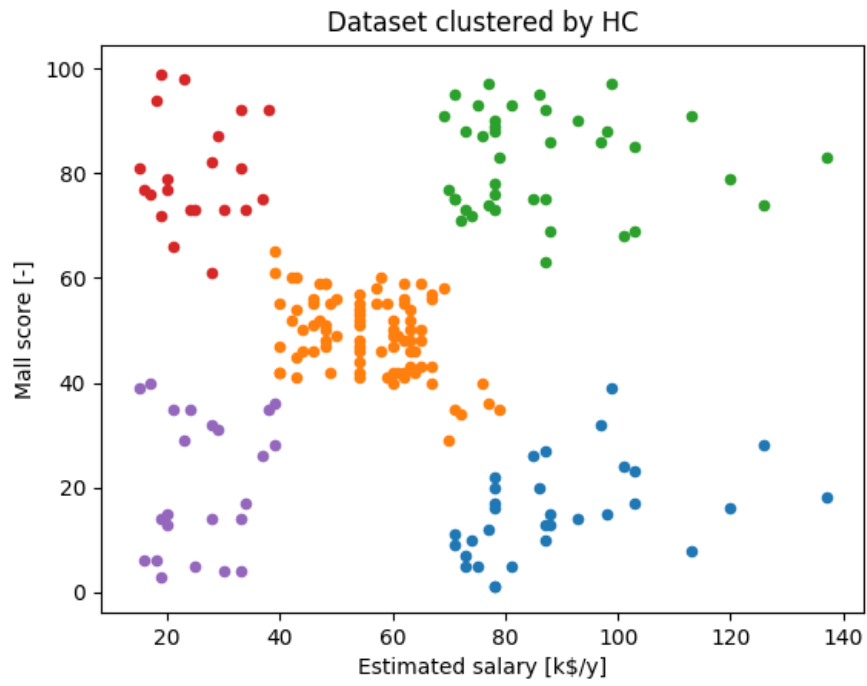
To choose the optimal  $k$  - the number of clusters, we use dendrograms. Dendrograms trace the history of points clustering and their distances (or dissimilarities). We choose the longest vertical line that can be split by horizontal line without interruption.



**Fig. 23.** Dendrogram - scipy library in Python

### Code

linkage is the method that is trying to minimize the variance in each cluster.



**Fig. 24.** Result of hierarchical clustering in Python

## 5 Association rule learning

"People who bought ... also bought ..."

**Problem:** We have a dataset containing transactions in a mall. We want to find associative rules. Here is the sample of the dataset. One row is one basket bought in the mall.

shrimp	almonds	avocado	vegetables mix	green grapes	whole wheat flour	yams	cottage
cheese	energy drink	tomato juice	low fat yogurt	green tea	honey	salad	mineral

### 5.1 Apriori

The important variables. The examples are from the problem.

- Support (M1) =  $\frac{\text{transactions containing } M1}{\text{total transactions}}$
- Confidence (M1 → M2) =  $\frac{\text{transactions containing } M1 \text{ and } M2}{\text{transactions containing } M1}$
- Lift (M1 → M2) =  $\frac{\text{Confidence}(M1 \rightarrow M2)}{\text{Support}(M2)}$

The algorithm has following steps:

1. Select a minimum support and confidence.
2. Take all the subsets in transactions having higher support than minimum support.
3. Take all the rules of these subsets having higher confidence than minimum confidence.
4. Sort the rules by decreasing lift.

This code implements the Apriori algorithm in R.

```

1 # Apriori
2
3 # Data Preprocessing
4 #install.packages('arules')
5 library(arules)
6 dataset = read.csv('Market_Basket_Optimisation.csv', header = FALSE)
7 dataset = read.transactions('Market_Basket_Optimisation.csv', sep = ',', rm.duplicates = TRUE)
8 summary(dataset)
9 itemFrequencyPlot(dataset, topN = 10)
10
11 # Training Apriori on the dataset
12 # FIRST we set wanted support. In our case we want to aim for products
13 # which are bought at least 3 or 4 times a day. That means support = (4*7)/7501
14 # SECOND we set the confidence to default value (0.8) and see how many rules we find.
15 # if we don't find enough rules then we lessen the confidence
16 rules = apriori(data = dataset, parameter = list(support = 0.003, confidence = 0.2))
17
18 # Visualising the results
19 inspect(sort(rules, by = 'lift')[1:10])

```

```

> summary(dataset)
transactions as itemMatrix in sparse format with
7501 rows (elements/itemsets/transactions) and
119 columns (items) and a density of 0.03288973

most frequent items:
mineral water      eggs      spaghetti french fries      chocolate      (Other)
      1788      1348      1306      1282      1229      22405

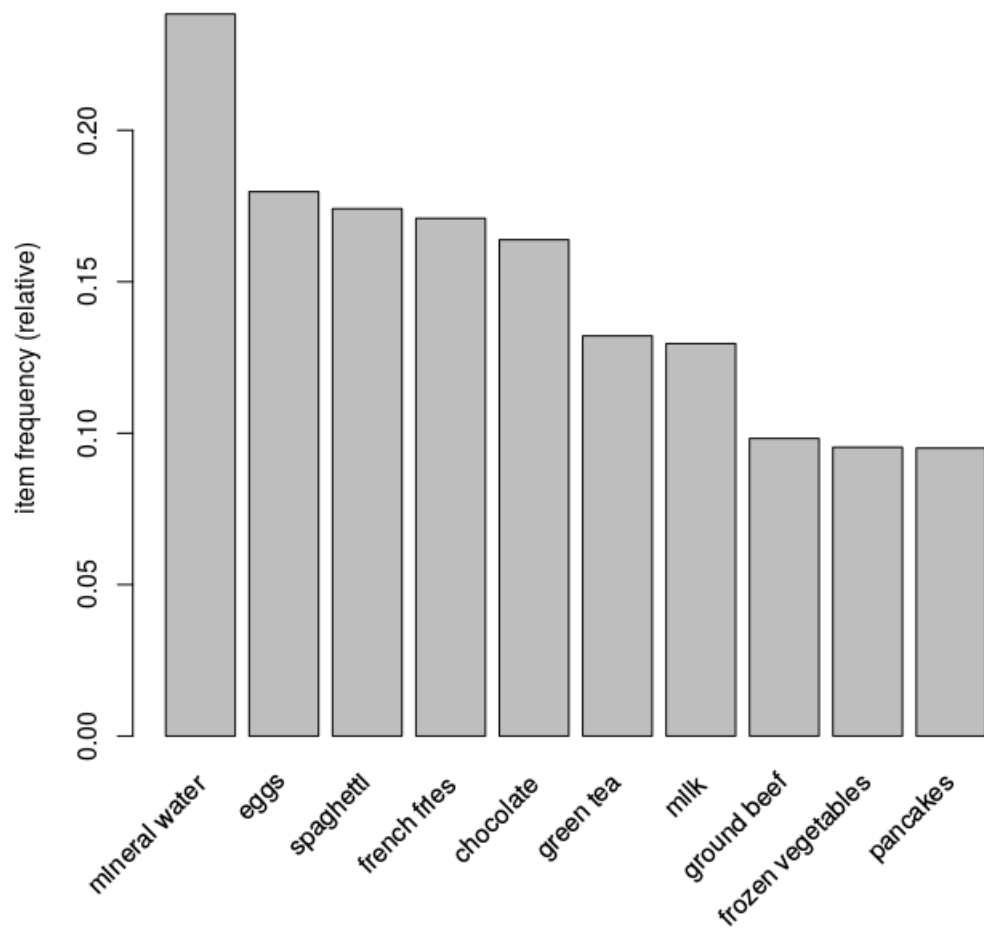
element (itemset/transaction) length distribution:
sizes
  1   2   3   4   5   6   7   8   9  10  11  12  13  14  15  16  18
1754 1358 1044 816 667 493 391 324 259 139 102 67 40 22 17 4 1
19 20
 2   1

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
      1.000  2.000   3.000   3.914  5.000  20.000

includes extended item information - examples:
      labels
1      almonds
2 antioxidant juice
3      asparagus
> |

```

**Fig. 25.** Summary of the R object. It is the matrix with a small amount of "ones" and a lot of "zeros". The number of rows equals the number of transactions and number of columns equals the cardinality of bought items set



**Fig. 26.** Visualisation of the most bought products. That means it shows their support

```

>
> inspect(sort(rules, by = 'lift')[1:10])
  lhs                                     rhs      support  confidence lift
[1] {mineral water,whole wheat pasta} => {olive oil} 0.003866151 0.4027778 6.115863
[2] {frozen vegetables,milk,mineral water} => {soup} 0.003066258 0.2771084 5.484407
[3] {fromage blanc} => {honey} 0.003332889 0.2450980 5.164271
[4] {spaghetti,tomato sauce} => {ground beef} 0.003066258 0.4893617 4.980600
[5] {light cream} => {chicken} 0.004532729 0.2905983 4.843951
[6] {pasta} => {escalope} 0.005865885 0.3728814 4.700812
[7] {french fries,herb & pepper} => {ground beef} 0.003199573 0.4615385 4.697422
[8] {cereals,spaghetti} => {ground beef} 0.003066258 0.4600000 4.681764
[9] {frozen vegetables,mineral water,soup} => {milk} 0.003066258 0.6052632 4.670863
[10] {french fries,ground beef} => {herb & pepper} 0.003199573 0.2307692 4.665768
> |

```

**Fig. 27.** Top 10 found rules in the dataset of bascet. Sorted by lift

## 5.2 Eclat

Eclat is similar to Apriori but much easier. We use only support.

The algorithm has following steps:

1. Select a minimum support.
2. Take all the subsets in transactions having higher support than minimum support.
3. Sort these subsets by decreasing support.

This code implements the Eclat algorithm in R. The only differences between this and Apriori is the function call and only support parameter.

```

1
2 # Training Eclat on the dataset
3 rules = eclat(data = dataset,
4               parameter = list(support = 0.003, minlen = 2))

```

## 6 Reinforcement learning

Reinforcement Learning is a branch of Machine Learning, also called Online Learning. It is used to solve interacting problems where the data observed up to time  $t$  is considered to decide which action to take at time  $t + 1$ . It is also used for Artificial Intelligence when training machines to perform tasks such as walking. Desired outcomes provide the AI with reward, undesired with punishment. Machines learn through trial and error.

### 6.1 Upper Confidence Bound (UCB)

### 6.2 Thompson sampling

## 7 Natural language processing

Natural Language Processing (or NLP) is applying Machine Learning models to text and language. Teaching machines to understand what is said in spoken and written word is the focus of Natural Language Processing. Whenever you dictate something into your iPhone / Android device that is then converted to text, that's an NLP algorithm in action.

You can also use NLP on a text review to predict if the review is a good one or a bad one. You can use NLP on an article to predict some categories of the articles you are trying to segment. You can use NLP on a book to predict the genre of the book. And it can go further, you can use NLP to build a machine translator or a speech recognition system, and in that last example you use classification algorithms to classify language. Speaking of classification algorithms, most of NLP algorithms are classification models, and they include Logistic Regression, Naive Bayes, CART which is a model based on decision trees, Maximum Entropy again related to Decision Trees, Hidden Markov Models which are models based on Markov processes.

A very well-known model in NLP is the Bag of Words model. It is a model used to preprocess the texts to classify before fitting the classification algorithms on the observations containing the texts.

In this part, you will understand and learn how to:

- Clean texts to prepare them for the Machine Learning models,
- Create a Bag of Words model,
- Apply Machine Learning models onto this Bag of Worlds model.



## 8 Deep learning

Deep Learning is the most exciting and powerful branch of Machine Learning. Deep Learning models can be used for a variety of complex tasks:

- Artificial Neural Networks for Regression and Classification
- Convolutional Neural Networks for Computer Vision
- Recurrent Neural Networks for Time Series Analysis
- Self Organizing Maps for Feature Extraction
- Deep Boltzmann Machines for Recommendation Systems
- Auto Encoders for Recommendation Systems

In this part, you will understand and learn how to implement the following Deep Learning models:

1. Artificial Neural Networks for a Business Problem
2. Convolutional Neural Networks for a Computer Vision task

### 8.1 Artificial Neural Network

### 8.2 Convolutional Neural Network

## **9   Dimensionality reduction**

### **9.1   Principal Component Analysis (PCA)**

### **9.2   Linear Discriminant Analysis (LDA)**

### **9.3   Kernel PCA**

## **10 Model selection and boosting**

### **10.1 Model selection**

### **10.2 XGBoost**